



TENSORFLOW 2.X CRASH COURSE

For AOS551 Fall 2021 Deep Learning in Geophysical Fluid Dynamics

Ming-Ruey (Ray) Chou

LEARNING RESOURCES




Official Site

Guide and Tutorial

Great explanation, detailed and thorough

Simply too much, hard to get started (need hours of reading)

Focus on traditional deep learning (CV, NLP, RL), which may not suit our needs



Very hands-on, thus lack some key details.

Suitable as a starter, but after which one still needs the official site.

(Comments are for the 1st edition)

TOPICS COVERED

TENSOR

AND ITS OPERATION

NumPy Array, Tensor, Variable
Taking Gradients

BUILDING A MODEL

Raw Variables or Keras
Create/Save/Load Weights

INTERGRATE

Data loading
Training loop

TIPS & PITFALLS

Debugging tips
Optimization tips



TOPICS ***NOT*** COVERED

Eager execution/Graph Execution, or TensorFlow 2.X/TensorFlow 1.X

Stick to 2.X, until training speed bothers you

[Graphs and tf.function](#) explained by official site

tf.data.Dataset API

Use something simple, except large (> > 10G) data.

[Input pipelines](#) explained by official site

Certain Useful but Baffling Keras API

model.fit, mode.compile, Callbacks, Functional API, .etc.

For simple models and training steps, the time to learn these does not justify

[Keras](#) explained by official site

NUMPY ARRAY

ARRAY CREATION

```
In [1]: import numpy as np
```

```
In [2]: arr = np.array(  
...:   [[1, 2, 3], [4, 5, 6]]  
...:   )
```

```
In [3]: arr  
Out[3]:  
array([[1, 2, 3],  
       [4, 5, 6]])
```

```
In [4]: ones = np.ones((3, 4))
```

```
In [5]: ones  
Out[5]:  
array([[1., 1., 1., 1.],  
       [1., 1., 1., 1.],  
       [1., 1., 1., 1.]])
```

ARRAY PROPERTIES

```
In [4]: arr.dtype  
Out[4]: dtype('int64')
```

```
In [5]: ones.dtype  
Out[5]: dtype('float64')
```

```
In [6]: arr.shape  
Out[6]: (2, 3)
```

```
In [7]: ones.shape  
Out[7]: (3, 4)
```

```
In [8]: arr.size  
Out[8]: 6
```

```
In [9]: ones.size  
Out[9]: 12
```

NUMPY ARRAY

ARITHMETIC OPERATIONS

```
In [7]: arr * 10
Out[7]:
array([[10, 20, 30],
       [40, 50, 60]])

In [8]: arr ** 2
Out[8]:
array([[ 1,  4,  9],
       [16, 25, 36]])

In [9]: arr.dot(ones)
Out[9]:
array([[ 6.,  6.,  6.,  6.],
       [15., 15., 15., 15.]])
```

SLICING

```
In [3]: arr[0, 0]
Out[3]: 1.0

In [4]: arr[0, :]
Out[4]: array([1., 2., 3.])

In [5]: arr[0, 1:3]
Out[5]: array([2., 3.])

In [6]: arr[:, 1:3]
Out[6]:
array([[2., 3.],
       [5., 6.]])
```

TENSORFLOW TENSOR

TENSOR CREATION

```
In [1]: import tensorflow as tf
In [3]: ts = tf.constant(
...:   [[1, 2, 3], [4, 5, 6]]
...: )

In [4]: ts
Out[4]:
<tf.Tensor: shape=(2, 3), dtype=int32, numpy=
array([[1, 2, 3],
       [4, 5, 6]], dtype=int32)>

In [5]: ts_ones = tf.ones((3, 4), dtype=tf.float32)

In [6]: ts_ones
Out[6]:
<tf.Tensor: shape=(3, 4), dtype=float32, numpy=
array([[1., 1., 1., 1.],
       [1., 1., 1., 1.],
       [1., 1., 1., 1.]], dtype=float32)>
```

TENSOR PROPERTIES

```
In [7]: ts.dtype
Out[7]: tf.int32

In [8]: ts_ones.dtype
Out[8]: tf.float32

In [9]: ts.shape
Out[9]: TensorShape([2, 3])

In [10]: ts_ones.shape
Out[10]: TensorShape([3, 4])

In [11]: tf.size(ts)
Out[11]: <tf.Tensor: shape=(),
          dtype=int32, numpy=6>

In [12]: tf.size(ts_ones)
Out[12]: <tf.Tensor: shape=(),
          dtype=int32, numpy=12>
```

TENSORFLOW TENSOR

ARITHMETIC OPERATIONS

```
In [19]: ts * 10
Out[19]:
<tf.Tensor: shape=(2, 3), dtype=int32, numpy=
array([[10, 20, 30],
       [40, 50, 60]], dtype=int32)>
```

```
In [20]: ts ** 2
Out[20]:
<tf.Tensor: shape=(2, 3), dtype=int32, numpy=
array([[ 1,  4,  9],
       [16, 25, 36]], dtype=int32)>
```

```
In [21]: tf.matmul(
...:     ts, tf.cast(ts_ones, dtype=tf.int32)
...: )
Out[21]:
<tf.Tensor: shape=(2, 4), dtype=int32, numpy=
array([[ 6,  6,  6,  6],
       [15, 15, 15, 15]], dtype=int32)>
```

SLICING

```
In [27]: ts[0, 0]
Out[27]: <tf.Tensor: shape=(),
dtype=int32, numpy=1>
```

```
In [28]: ts[0, :]
Out[28]: <tf.Tensor: shape=(3,),
dtype=int32, numpy=array([
1, 2, 3], dtype=int32)>
```

```
In [30]: ts[:, 1:3]
Out[30]:
<tf.Tensor: shape=(2, 2), dtype=int32, numpy=
array([[2, 3],
       [5, 6]], dtype=int32)>
```


Array – Tensor Operations

Automatic conversion (sometimes)

```
In [36]: arr.dot(ts_ones)
```

```
Out[36]:  
array([[ 6.,  6.,  6.,  6.],  
       [15., 15., 15., 15.]])
```

```
In [37]: arr * ts
```

```
Out[37]:  
<tf.Tensor: shape=(2, 3), dtype=int32, numpy=  
array([[ 1,  4,  9],  
       [16, 25, 36]], dtype=int32)>
```

Explicit conversion (via `.numpy()` / `tf.constant`)

```
In [38]: ts.numpy( )
```

```
Out[38]:  
array([[1, 2, 3],  
       [4, 5, 6]], dtype=int32)
```

```
In [39]: tf.constant(arr, dtype=tf.int32)
```

```
Out[39]:  
<tf.Tensor: shape=(2, 3), dtype=int32, numpy=  
array([[1, 2, 3],  
       [4, 5, 6]], dtype=int32)>
```

TENSOR/NUMPY
ESSENTIALS

Tensor/NumPy API Similarity

One could expect, at least for the basic NumPy ndarray operations, TensorFlow provides similar syntax or functionality

However, aiming for parallel computing makes some fundamental differences between the two (similar operations still included, but wonkier and less computationally efficient)

We cover some important differences between them in next slides

Array-Like, but *Strict Type*

Array, Automatic Type Conversion:

```
In [28]: arr * ones[0:2, 0:3]
Out[28]:
array([[1., 2., 3.],
       [4., 5., 6.]])
```

Note: dtype of arr is int64, ones is float64

Tensor Operation is Strict Type

```
In [29]: ts * ts_ones[0:2, 0:3]
```

InvalidArgumentError:

cannot compute Mul as input #1(zero-based) was expected to be an int32 tensor but is a float tensor [Op:Mul]

TENSOR/NUMPY ESSENTIALS

Need to Cast Datatype before Operations:

```
In [30]: tf.cast(ts, dtype=tf.float32) * ts_ones[0:2, 0:3]
Out[30]:
<tf.Tensor: shape=(2, 3), dtype=float32, numpy=
array([[1., 2., 3.],
       [4., 5., 6.]], dtype=float32)>
```

but Hard for Advanced Slicing

Advanced Slicing for Array:

```
In [16]: arr_indices = np.array(
...:     [[0, 1], [1, 0]])
```

```
In [17]: arr[arr_indices]
```

```
Out[17]:
array([[1, 2, 3],
       [4, 5, 6]],

      [[4, 5, 6],
       [1, 2, 3]])
```

Results in Error for Tensor

```
In [19]: ts_indices = tf.constant(
...:     [[0, 1], [1, 0]])
```

```
In [20]: ts[ts_indices]
InvalidArgumentError
```

TENSOR/NUMPY
ESSENTIALS

Can achieve similar operation in TensorFlow,
but cumbersome.

but Immutable for 'normal' Tensor

Change Value(s) in Array

```
In [31]: arr[0, 0] = 10
```

```
In [32]: arr
```

```
Out[32]:  
array([[10,  2,  3],  
       [ 4,  5,  6]])
```

Results in Error for Tensor

```
In [33]: ts[0, 0] = 10
```

TypeError:

'tensorflow.python.framework.ops.EagerTensor'
object does not support item assignment

TENSOR/NUMPY ESSENTIALS

Tensor in general is immutable.

The only way to change values in existing
Tensor is to use **Variable**.

VARIABLE, Tensor whose values are modifiable

IS A TENSOR

```
In [3]: var = tf.Variable(  
...: [[1, 2, 3], [4, 5, 6]])
```

```
In [4]: var * 10
```

```
Out[4]:  
<tf.Tensor: shape=(2, 3), dtype=int32, numpy=  
array([[10, 20, 30],  
       [40, 50, 60]]), dtype=int32)>
```

```
In [5]: var[0, :2]
```

```
Out[5]: <tf.Tensor: shape=(2,), dtype=int32,  
        numpy=array([1, 2], dtype=int32)>
```

WITH METHOD TO CHANGE ITS VALUES

```
In [8]: var.assign([[7, 8, 9], [10, 11, 12]])
```

```
Out[8]:  
<tf.Variable 'UnreadVariable' shape=(2, 3) dt  
array([[ 7,  8,  9],  
       [10, 11, 12]], dtype=int32)>
```

```
In [9]: var
```

```
Out[9]:  
<tf.Variable 'Variable:0' shape=(2, 3) dtype=  
array([[ 7,  8,  9],  
       [10, 11, 12]], dtype=int32)>
```

REMARK - Model weights are Variables

REMARKS

KISS (***Keep it simple, stupid***)

- Do explicit conversion instead of automatic
- Stick to dumb simple operations
- Broadcasting is powerful but leads to obscure bugs. Keep dimension and shape as simple as possible (while considering batch dimension)

TENSOR/NUMPY ESSENTIALS

What is broadcast? General rule for NumPy array operation

Array (elementwise) Multiply:

```
A      (4d array):  8 x 1 x 6 x 1
B      (3d array):      7 x 1 x 5
Result (4d array):  8 x 7 x 6 x 5
```

From NumPy official site

TAKE GRADIENTS

```
x = tf.Variable(3.0)

with tf.GradientTape() as tape:
    y = x**2
```

Roughly equals:

```
tape = tf.GradientTape()
tape.__enter__()
```

```
# dy = 2x * dx
dy_dx = tape.gradient(y, x)
dy_dx.numpy()
```

>>> will get 6.0

The GradientTape is the fundamental tool in TensorFlow to do automatic differentiation (i.e., calculate gradient)

To understand how automatic differentiation works, see:

Baydin, A., et al. (2017). Automatic Differentiation in Machine Learning: A Survey. J. Mach. Learn. Res., 18(1), 5595–5637.

TAKE GRADIENTS

```
w = tf.Variable(tf.random.normal((3, 2)), name='w')
b = tf.Variable(tf.zeros(2, dtype=tf.float32), name='b')
x = [[1., 2., 3.]]
```

```
with tf.GradientTape(persistent=True) as tape:
```

```
    y = x @ w + b
```

```
    loss = tf.reduce_mean(y**2)
```

Parameter to control if the tape is used
one-time-only

Block for tape to trace the gradient

Tape traces operations within the block only

```
[dl_dw, dl_db] = tape.gradient(loss, [w, b])
```

```
[dy_dw, dy_db] = tape.gradient(y, [w, b])
```

Use the same tape to get gradient second
time won't work without being persistent

TAKE GRADIENTS

```
# A trainable variable
x0 = tf.Variable(3.0, name='x0')
# Not trainable
x1 = tf.Variable(3.0, name='x1', trainable=False)
# Not a Variable: A variable + tensor returns a tensor.
x2 = tf.Variable(2.0, name='x2') + 1.0
# Not a variable
x3 = tf.constant(3.0, name='x3')

with tf.GradientTape() as tape:
    y = (x0**2) + (x1**2) + (x2**2)

grad = tape.gradient(y, [x0, x1, x2, x3])

for g in grad:
    print(g)
```

Yield proper gradients only for x0

```
tf.Tensor(6.0, shape=(), dtype=float32)
None
None
None
```

GradientTape traces only
trainable Variables by default

Note:

x1 is not-trainable

x2 is a Tensor! (The addition +1.0 returns a Tensor)

TAKE GRADIENTS

Explicit stated which Tensor/Variable you want to get derivatives

```
x = tf.constant(3.0)
with tf.GradientTape() as tape:
    tape.watch(x)
    y = x**2

# dy = 2x * dx
dy_dx = tape.gradient(y, x)
print(dy_dx.numpy())
```

>>> will get 6.0

All the parameters along the way are traced, so we can get derivative w.r.t to intermediate result

```
x = tf.constant(3.0)

with tf.GradientTape() as tape:
    tape.watch(x)
    y = x * x
    z = y * y

# Use the tape to compute the gradient of z with
# intermediate value y.
# dz_dy = 2 * y and y = x ** 2 = 9
print(tape.gradient(z, y).numpy())
```

TAKE GRADIENT AND UPDATE

Simple Gradient Descent example:

manually calculating the gradients and updating the values

```
with tf.GradientTape() as t:
    # Trainable variables are automatically tracked by GradientTape
    current_loss = loss(y, model(x))

    # Use GradientTape to calculate the gradients with respect to W and b
    dw, db = t.gradient(current_loss, [model.w, model.b])

    # Subtract the gradient scaled by the learning rate
    model.w.assign_sub(learning_rate * dw)
    model.b.assign_sub(learning_rate * db)
```

TAKE GRADIENT AND UPDATE, USING BUILT-IN OPTIMIZERS

apply_gradients wraps all the calls required for updating weights, also the calculation for learning rate .etc.

```
adam = tf.keras.optimizers.Adam(learning_rate=0.01)

with tf.GradientTape() as tape:
    losses = loss(y, model(x))

gradients = tape.gradient(losses, model.trainable_variables)
adam.apply_gradients(zip(gradients, model.trainable_variables))
```

KERAS INTERFACE

It wraps the Tensor/Variable as realization of high level concepts such as layers, weight initialization, model, optimizer .etc

Working in high-level concept is time saving and less error-prone

LAYERS & MODELS IN KERAS

CREATE A LAYER

```
layer = Dense(50, activation="tanh",  
             dtype=tf.float32,  
             kernel_initializer="glorot_normal")
```

It's a fully-connected layer of 50 nodes with tangent hyperbolic activation

```
input_layer = InputLayer(input_shape=(1,))
```

InputLayer is a specialized layer without any weights & activation

It just determine the dimension of the input data (1D input in this case)

STACK LAYERS TO MODEL

```
from tensorflow.keras.layers import Dense, InputLayer  
  
model = keras.Sequential()  
model.add(InputLayer(input_shape=(1,)))  
model.add(  
    Dense(50, activation="tanh",  
         dtype=tf.float32,  
         kernel_initializer="glorot_normal"))  
model.add(  
    Dense(1, activation="tanh",  
         dtype=tf.float32,  
         kernel_initializer="glorot_normal"))
```

REMARK – Check the `tf.keras.layers` module most of the commonly used operations are already there

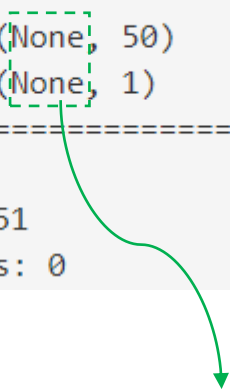
LAYERS & MODELS IN KERAS

THE STRUCTURE OF MODEL

```
model.summary()
```

display the structure of the neural network

Layer (type)	Output Shape	Param #
=====		
dense (Dense)	(None, 50)	100
dense_1 (Dense)	(None, 1)	51
=====		
Total params: 151		
Trainable params: 151		
Non-trainable params: 0		



REMARK - the **batch dimension** is automatically taking care of, which reduces much trouble.

FORWARD-PASS

```
# shape of x1 is (1, 1)
# model expect input has batch dimension
#
# y1 is a Tensor with shape (1, 1)
x1 = tf.constant([[0.0]])
y1 = model(x1)
```

Model expects input data shape is (batch, 1)

```
# shape of x2 is (100, 1)
# y2 is a Tensor with shape (100, 1)
x2 = tf.ones((100, 1))
y2 = model(x2)
```

Model outputs shape is also (batch, 1)

LAYERS & MODELS IN KERAS

FORWARD-PASS 2D EXAMPLE

```
model = keras.Sequential()
model.add(InputLayer(input_shape=(2,)))
model.add(Dense(50))
model.add(Dense(3))

model.summary()

# output is a Tensor with shape (100, 3)
points = tf.ones((100, 2))
outputs = model(points)

# u, v, h are with shape (100,)
velocity_u = outputs[:, 0]
velocity_v = outputs[:, 1]
hardness_h = outputs[:, 2]
```

Specify your input as 2D, so the model expects the input shape to be (batch, 2)

Layer (type)	Output Shape	Param #
=====		
dense (Dense)	(None, 50)	150
dense_1 (Dense)	(None, 3)	153
=====		
Total params: 303		
Trainable params: 303		
Non-trainable params: 0		

Model output shape is (batch, 3)

Slice the (stacked) model outputs to get the separate component of your physical system

RAW TENSOR OPERATIONS versus KERAS MODEL BUILDING

Using Raw Tensors/Variables

```
class MyNet:

    def __init__(self):
        layers = [1, 50, 50, 50, 1]

        # initialize weights and biases
        initializer = tf.keras.initializers.GlorotNormal()
        self.weights = []
        self.biases = []
        for l in range(len(layers) - 1):
            W = initializer(shape=(layers[l], layers[l + 1]))
            b = tf.Variable(tf.zeros([1, layers[l + 1]], dtype=tf.float32))
            self.weights.append(W)
            self.biases.append(b)
        self.trainable_variables = self.weights + self.biases

    def forward(self, x):
        H = x
        for l in range(len(self.weights) - 1):
            W = self.weights[l]
            b = self.biases[l]
            H = tf.tanh(tf.add(tf.matmul(H, W), b))
        W = self.weights[-1]
        b = self.biases[-1]
        Y = tf.add(tf.matmul(H, W), b)
        return Y
```

Using Keras

```
model = keras.Sequential()
model.add(InputLayer(input_shape=(1,)))
model.add(
    Dense(50, activation="tanh",
          dtype=tf.float32,
          kernel_initializer="glorot_normal"))
model.add(
    Dense(50, activation="tanh",
          dtype=tf.float32,
          kernel_initializer="glorot_normal"))
model.add(
    Dense(50, activation="tanh",
          dtype=tf.float32,
          kernel_initializer="glorot_normal"))
model.add(
    Dense(1, activation="tanh",
          dtype=tf.float32,
          kernel_initializer="glorot_normal"))
model.summary()
```

USING KERAS MODEL

WEIGHTS CREATED AUTOMATICALLY

```
In [12]: model.layers
```

```
Out[12]:  
[<keras.layers.core.Dense at 0x7fe916c38ac8>,  
 <keras.layers.core.Dense at 0x7fe916c38dd8>]
```

```
In [13]: model.layers[0].weights
```

```
Out[13]:  
[<tf.Variable 'dense_2/kernel:0' shape=(2, 50) dtype=float32, numpy=  
 array([[ -0.27954745, -0.08841521,  0.15327993,  0.20243871, -0.2656639  
        ... dtype=float32)>,  
 <tf.Variable 'dense_2/bias:0' shape=(50,) dtype=float32, numpy=  
 array([0., 0., 0., 0., 0., 0., 0., 0., 0., 0., 0., 0., 0., 0., 0.,  
        ... dtype=float32)>]
```

Get a list of all trainable weights via `model.trainable_variables`,
which is useful for updating weights

```
In [4]: model.trainable_variables
```

USING KERAS MODEL

UPDATE VARIABLES

```
adam = tf.keras.optimizers.Adam(learning_rate=0.01)

with tf.GradientTape() as tape:
    losses = loss(y, model(x))

gradients = tape.gradient(losses, model.trainable_variables)
adam.apply_gradients(zip(gradients, model.trainable_variables))
```

SAVE/LOAD WEIGHTS

```
model.save_weights("my_model.h5")

# You have to recreate a model with EXACT same structure
model.load_weights("my_model.h5")
```

CAVEAT – load_weights with different model structure may silently fail, without any apparent warning/exception

OTHER CUSTOMIZABLE

CUSTOM LAYERS

```
class MyDenseLayer(tf.keras.layers.Layer):
    def __init__(self, num_outputs):
        super(MyDenseLayer, self).__init__()
        self.num_outputs = num_outputs

    def build(self, input_shape):
        self.kernel = self.add_weight("kernel",
                                       shape=[int(input_shape[-1]),
                                              self.num_outputs])

    def call(self, inputs):
        return tf.matmul(inputs, self.kernel)
```

CUSTOM INITIALIZER

```
class ExampleRandomNormal(tf.keras.initializers.Initializer):

    def __init__(self, mean, stddev):
        self.mean = mean
        self.stddev = stddev

    def __call__(self, shape, dtype=None, **kwargs):
        return tf.random.normal(
            shape, mean=self.mean, stddev=self.stddev, dtype=dtype)

    def get_config(self): # To support serialization
        return {"mean": self.mean, "stddev": self.stddev}
```

EVEN CUSTOM OPTIMIZER! ...though harder

REMARKS - usually not needed, but can be useful sometimes

Also, checkout [this wonderful colab demonstration](https://colab.research.google.com/drive/17u-pRZJnKN0gO5XZmq8n5A2bKGrfKEUg#scrollTo=UHOOIxcQ9GI)

<https://colab.research.google.com/drive/17u-pRZJnKN0gO5XZmq8n5A2bKGrfKEUg#scrollTo=UHOOIxcQ9GI>

A series of thin, light-brown lines forming various overlapping triangles and polygons, creating an abstract geometric pattern in the background.

INTERGRATE EVERYTHING

Pseudo code demonstrates the fundamental elements of a training script

```
x_data, y_data = ... # get your data  
model = ... # create model  
opt = ... # Use built-in optimizer
```

```
for iteration in range(...):  
    with tf.GradientTape() as tp:  
        prediction = model(x_data)  
        loss = MSE(y_data, prediction)
```

```
weights = model.trainable_variabels  
grads = tp.gradient(loss, weights)  
opt.apply_gradients(zip(grads, weights))
```

```
def prepare_data():
    x_data, y_data = LOAD("data.mat")
    x_data /= XSCALE
    y_data /= YSCALE
    x_data = x_data[..., np.newaxis]
    y_data = y_data[..., np.newaxis]
    x_tensor = tf.constant(x_data)
    y_tensor = tf.constant(y_data)
    return x_tensor, y_tensor
```

DATA LOADING

- Load
- Normalize
- Correct the Shape (First dimension is batch!)
- (Optional) Sample Data in Loop

When Data Fits Into Memory

Load entire data into memory before the training loop. Can choose to sample a subset at each iteration

```
random_shuffle(x_data, y_data)
for iteration in range(...):
    start = (iteration % 100) * 1000
    end = start + 999
    x_sampled = x_data[start:end, :]
    y_sampled = y_data[start:end, :]
    ...
```

When Dataset Is Large

Harder! Create an iterator over the model, fetch the data on the fly. Checkout [input pipelines](#) explained by official site

```
x_data, y_data = ... # get your data
model = ... # create model
opt = ... # Use built-in optimizer
```

```
for iteration in range(...):
    with tf.GradientTape() as tp:
        prediction = model(x_data)
        loss = MSE(y_data, prediction)

    weights = model.trainable_variables
    grads = tp.gradient(loss, weights)
    opt.apply_gradients(zip(grads, weights))
```

INTERGRADE EVERYTHING

```
def equation_loss(x, model):
    with tf.GradientTape(persistent=True) as tp:
        tp.watch(x)
        outputs = model(x)
        u = outputs[..., 0]
        v = outputs[..., 1]

    u_x = tp.gradient(u, x)
    v_x = tp.gradient(v, x)

    # Residues of Governing Equations
    equation = ALPHA * u_x - BETA * v_x
    return tf.reduce_mean(equation ** 2)
```

```
x_data, y_data = ... # get your data
model = ... # create model
opt = ... # Use built-in optimizer
```

```
for iteration in range(...):
    with tf.GradientTape() as tp:
        prediction = model(x_data)
        loss = MSE(y_data, prediction)

    weights = model.trainable_variables
    grads = tp.gradient(loss, weights)
    opt.apply_gradients(zip(grads, weights))
```

INTERGRADE EVERYTHING

OPTIMIZATION STEPS

- Calculate Loss
- Calculate Gradients
- Update Variables
- Display/Save for Monitoring

Monitor Is A Must

Loss values on train/validation set for every n iterations

Include any relevant metrics

Store Intermediate Models Is Helpful

Robust against unexpected interruption

Easiest way to select model based on validation set

```
for iteration in range(...):
    ...

    if iteration % 100 == 0:
        print(loss.numpy())
        model.save(f"at_it={iteration}.h5")
```


General Principle of Coding:

Keep Steps Small And Separated

- ❖ Debugging effort determined (almost solely) by suspected area
- ❖ Splitting code into different files helps
- ❖ Create small and simple functions you are confident with, and rely on them
- ❖ Able to revert to the previous working version (version control tool is helpful)

MODEL TRAINING PITFALLS

Use Simple .py file Instead of Notebook

Notebook remembers variable value even the cell is deleted or changed

The execution order may affect the result

Not suitable as scheduled job on cluster

Split Training and Plotting Code

Do NOT use an end-to-end code file for everything. Split, split, split!

Training and store the model using a stand-alone .py file

Load the stored model in another file to examination the result

Start with Simplified Case as Sanity Check

Deep learning, especially the optimization, is hard to reason about

Ground project on simplified situation which we have clear expected result

Develop toward complex and long-running

Checklist When Model Performs Unreasonably Poor

- Have some simplified case that works as sanity check
- Does the model do the right calculation?
 - Careful on the numerical range of activation function e.g., tanh only outputs in $[-1, 1]$
 - Check the shape of input and output for each layer
- Is the long-term trend of loss curve decreasing?
 - Does the model weights change at all?
 - Try different learning rates
 - Try different initializers
- Are the loss values sane?
- Learn from the Model Prediction
 - In which domain, it performs poorly?
 - Is there any patterns in the prediction?

MODEL TRAINING PITFALLS

A series of thin, light brown lines forming an abstract geometric pattern in the top left corner of the slide. The lines intersect to create various triangular and polygonal shapes.

ENJOY, GOOD LUCK

REACH FOR HELP

Ming-Ruey (Ray) Chou, mc4536@princeton.edu,

Office hour: Tuesday 13:00 – 14:00

- AMA via email/canvas
- Other slots are available via appointment